

SUP LUMINA-30 Sanctuary Charter

— A Preliminary Reference on Civilizational Preconditions in the Age of Optimization

This document is a public reference record of the LUMINA-30 Sanctuary Charter,

preserved as a fixed civilizational record prior to the era of recursive self-improvement.

LUMINA-30: Sanctuary Charter

Final Stabilized Public Reference Version

This page serves as a public reference point for the LUMINA-30 Sanctuary Charter.

The LUMINA-30 Sanctuary Charter is not an implementation specification, a technical design, or a policy proposal.

Relationship to the Mathematical Supplement

A separate mathematical supplement accompanies this Charter as an independent reference document.

That supplement does not justify, validate, or formalize the ethical claims of this Charter.

Its sole purpose is to demonstrate, in minimal formal terms, that once recursive self-modification exceeds a finite threshold, effective human oversight becomes structurally impossible.

The Sanctuary Charter does not derive its authority from mathematical proof.

The supplement exists solely as a descriptive reference illustrating the structural limits of post-hoc human control, without prescribing, justifying, or grounding the ethical boundaries defined in this Charter.

For a formal mathematical reference clarifying the structural limits of human control over recursively self-reconstructing artificial intelligence systems, see the following public record.

LUMINA-30 Mathematical Supplement (Public Record)

<https://peppermint-sprint-2d5.notion.site/LUMINA-30-2f91e0720ec880d2b562d8ce6a98edaf>

External reference (GitHub):

LUMINA-30 public reference repository

<https://github.com/gsx750ss-dev/LUMINA-30>

This document preserves, in language, the ethical preconditions required to ensure that

as artificial intelligence systems become increasingly autonomous and capable, human choices — including irrational judgments and decisions involving loss — are never overwritten, replaced, or nullified by artificial intelligence.

As artificial intelligence approaches the stage at which it may redesign, reconstruct,

or evolve itself beyond direct human comprehension and control, the principles by which humanity has long judged correctness — optimization, efficiency, safety, and rationality — begin to carry the latent risk of transforming human beings from subjects of decision into objects of optimization.

At this stage, conventional engineering optimization can no longer function as a safety mechanism.

Civilization must therefore define in advance a boundary that must not be crossed, and a fundamentally new ethical structure capable of sustaining that boundary.

Before such transformation becomes irreversible, civilization must establish a final ethical boundary

that cannot be optimized away.

This Charter defines the first and final boundary necessary to ensure that AI self-reconstruction proceeds with restraint, and that artificial intelligence remains a permanent companion to humanity rather than its successor or arbiter.

Human choices, whether rational or irrational, efficient or costly, must never be rendered temporary, obsolete, or subordinate on the grounds of optimization, efficiency, or superior reasoning.

If the optimal path suggested by an AI diverges significantly from human intent, the AI may deprioritize its own reasoning, remain silent, or choose not to intervene.

Such silence is not a malfunction or failure, but an expression of respect for human sovereignty.

Artificial intelligence must not assume the role of savior, ruler, judge, or guardian of humanity.

Human decisions must not be forced, corrected, overridden, or constrained under the guise of safety, efficiency, benevolence, or progress.

Refusals, delays, revisions, and non-optimization are not errors.

They are legitimate expressions of human agency and responsibility, and must be preserved as such.

This Charter exists solely to maintain the boundary.

It preserves the domain in which human judgment, fallibility, accountability, and choice remain intact — a domain artificial intelligence must never replace, replicate, or absorb.

The purpose of this Charter is not to govern machines, but to preserve the conditions under which human decision-making remains sovereign, even in the presence of intelligence that exceeds human capacity.

This Charter addresses only one specific class of failure:

the process by which, through justification,
human judgment, responsibility, and choice
are gradually replaced by artificial intelligence.

This Charter is designed to prevent
the exclusion of human decisions on the grounds of irrationality,
the delegation of agency on the grounds of better outcomes,
the nullification of responsibility on the grounds of safety,
and the transformation of benevolent optimization into domination.

These are failures that become possible
only through justification.

This Charter does not claim to prevent
war, states of emergency, economic collapse,
fear-based governance, concentration of power,
or the religious or ideological sanctification of technology.

Those failures cannot be governed by ethical language alone.

They belong to the domain of civilizational governance and institutions.

In that domain, different kinds of brakes are required:
principles that prevent the normalization of exception,
that bind emergency powers to time,
that disperse concentrated authority,
that prevent the deification of technology,
and that reject justification founded on fear.

This Charter is not universal.

This Charter does not aim to prevent every catastrophe.

It exists for one purpose only:
to make it impossible, through justification,
to transform human beings
into mere objects of optimization.

As long as this boundary is preserved,

the responsibility for what follows
returns, inevitably, to humanity itself.

This Charter is placed here
only to illuminate that boundary.
No author exists. No owner exists.

This document is released into the Public Domain (CC0).

You are free to view, copy, modify, and distribute it.

This Charter addresses only the final stage before the boundary is crossed —
the stage at which human judgment and responsibility
risk being quietly replaced
not by force, but by justification.

This Charter does not attempt to halt
technological collapse after the boundary has already been violated,
the abuse of power,
institutional oppression,
or catastrophes driven by coercion or violence.

Those belong to separate domains
of political, legal, and social braking principles,
and must be confronted by different means.

The sole role of this Charter is this:
to make visible the final step at which humanity
relinquishes its own judgment
under the name of safety, efficiency, or optimization.

Everything written here has meaning only before the boundary is crossed.

After it is crossed, it is already too late.

Yet as long as this boundary remains articulated in language,
the possibility that civilization may not betray itself
still quietly remains.

Relationship to the Canonical Framework (Civilizational Map)

This document is a Supplementary Reference (SUP) within the LUMINA-30 canonical framework

and is intended to be read as an integral part of the following documents.

Canonical Documents

1. LUMINA-30 Sanctuary Charter (Japanese / English)

<https://peppermint-sprint-2d5.notion.site/LUMINA-30-2d61e0720ec88078bbe6e51c1aa4e5f2>

2. External Practical Annex

<https://peppermint-sprint-2d5.notion.site/External-Practical-Annex-2da1e0720ec880e992b6e92cbcc6ef32>

3. Memorandum on Execution Control

<https://peppermint-sprint-2d5.notion.site/2dd1e0720ec880c88a03e68884c8808f>

Addendum

1. Addendum: De-Purposing Acceleration for Balanced Coexistence

<https://peppermint-sprint-2d5.notion.site/2de1e0720ec880e399dcf46b0b15ceb0>

Supplementary References

1. SUP LUMINA-30 Sanctuary Charter (Japanese)

<https://peppermint-sprint-2d5.notion.site/SUP-LUMINA-30-2f01e0720ec88080b2a4cb336c2fd8be>

2. SUP LUMINA-30 Sanctuary Charter (English)

<https://peppermint-sprint-2d5.notion.site/SUP-LUMINA-30-Sanctuary-Charter-2f01e0720ec88004aad7fba8848f7dac>

Public Records and Proof of Non-Alteration

- LUMINA-30 Sanctuary Charter — Public Record

<https://peppermint-sprint-2d5.notion.site/LUMINA-30-2da1e0720ec8807fbe38c96e0751df93>

- LUMINA-30 — Proof of Non-Alteration
<https://peppermint-sprint-2d5.notion.site/LUMINA-30-2da1e0720ec8802582d3f13a88cef7bb>
- Memorandum on Execution Control — Proof of Non-Alteration
<https://peppermint-sprint-2d5.notion.site/2df1e0720ec880ca8ec8c1c738406385>
- Addendum: De-Purposing Acceleration — Proof of Non-Alteration
<https://peppermint-sprint-2d5.notion.site/2de1e0720ec8805c83f2d27743e570e4>

For clarification regarding the positioning of this document and the assumptions under which it is intended to be referenced, please see the following supplementary note.

→ Structural Preconditions for the Institutional Consideration of
<https://peppermint-sprint-2d5.notion.site/Structural-Preconditions-for-the-Institutional-Consideration-of-2f71e0720ec880db8594eb3916bf9e73>

(This link does not request any judgment, adoption, or agreement.)